


# Airline Database



Based on the Airline database, write a SQL query to calculate how many reservations were made with the JetBlue flights for a cost higher than the average cost of all JetBlue reservations?

Note that the cost of a reservation is the sum of the price of all tickets in that reservation. Only list the total number of reservations. Rename the calculated field NumReservations>AvgForJB.

Step by step solution:

-- average per reservation

```

select r.rid, avg(t.price) as avg -- (A)
from airline a, flight f, reservation r, ticket_info t
where name='JetBlue'
and a.aid=f.aid and f.fid=r.fid and r.rid=t.rid
group by r.rid;
    
```

-- total average

```

select avg(t.price) as avg -- (B)
from airline a, flight f, reservation r, ticket_info t
where name='JetBlue'
and a.aid=f.aid and f.fid=r.fid and r.rid=t.rid
group by a.aid
    
```

```

select count(avg_per_reservation.rid) as 'NumReservations>AvgForJB'
from
    
```

```

    (select r.rid, avg(t.price) as avg -- (A)
    from airline a, flight f, reservation r, ticket_info t
    where name='JetBlue'
    and a.aid=f.aid and f.fid=r.fid and r.rid=t.rid
    group by r.rid) avg_per_reservation,
    
```


```

    (select avg(t.price) as avg -- (B)
    from airline a, flight f, reservation r, ticket_info t
    where name='JetBlue'
    and a.aid=f.aid and f.fid=r.fid and r.rid=t.rid
    group by a.aid) avg_total
    
```

```

where avg_per_reservation.avg > avg_total.avg;
    
```

# Stock Database



Based on the Stock database, write a SQL query to calculate the number of cities with 3 or more customers having completed a transaction. Only list the total number of cities, rename the calculated field NumCitiesLotsOfCustomers.

Step by step solution

-- get all cities

```
select z.city, z.zip
```

```
from zips z;
```

-- get all customers and zip

```
select z.city, z.zip, c.last_name
```

```
from zips z, customers c
```

```
where z.zip=c.zip;
```

-- count customers per zip

```
select z.city, z.zip, count(c.cid)
```

```
from zips z, customers c
```

```
where z.zip=c.zip
```

```
group by z.city, z.zip;
```

-- count customers per zip that have a portfolio

```
select z.city, z.zip, count(c.cid) -- (A)
```

```
from zips z, customers c, portfolios p
```

```
where z.zip=c.zip
```

```
and p.cid=c.cid
```

```
group by z.city, z.zip;
```

-- count cities that have 3 or more customers with portfolios

```
select count(customer_count.city)
```

```
from zips z,
```

```
    (select z.city, z.zip, count(c.cid) as count -- (A)
```

```
    from zips z, customers c, portfolios p
```


```
    where z.zip=c.zip
```

```
    and p.cid=c.cid
```

```
    group by z.city, z.zip) customer_count
```

```
where customer_count.count >= 3
```

```
and z.zip=customer_count.zip
```



## Final Exam Topic Review

- Design
  - UML
    - Text → UML → Text
    - Aggregation / Composition
      - CREATE ... CASCADE DELETE - Composition
      - Otherwise - Aggregation
    - Inheritance - triangle on the abstract/base class
    - Reification → convert pure UML to something can be implemented as SQL
    - Mapping class / Association
    - 1 - many
    - many - many
  - Functional Dependencies
    - FD with non keys - bad
    - FD with keys - good
  - Normalization
    - 1NF - fields should be dependent on the key
    - 2NF - FD with part of key - compound keys - the whole key
    - 3NF - FD with some other non key field - nothing but the key
  - Relational Algebra
    - SELECT(TABLE, predicate) ⇒ Sub table with missing rows only the ones that match criteria
    - PROJECTIONS(TABLE, [field1, field2]) ⇒ same table with missing fields
    - PRODUCT(TABLE1, TABLE2) ⇒ all combinations of records in 1 and 2
    - JOIN(TABLE1, TABLE2, predicate)
- Relational Databases
  - SELECT
  - JOINS
  - GROUP BY
  - ~~HAVING~~
  - ~~INSERT~~
  - ~~UPDATE~~
  - ~~DELETE~~
- Programming the database
  - Stored procedures
    -
  - triggers
    - Execute for an event

- Before/After ⇒ Update/Delete/Insert
  - Can do Update/Delete/Insert
- ~~JDBC~~
- ORM
  - @Entity
  - @Table
  - @Id
  - @GeneratedValue
  - @OneToMany(mappedBy="")
  - @ManyToOne
  - @JsonIgnore
- NoSQL
  - MongoDB
    - show dbs - to list existing databases
    - use myDatabase - to create a database
    - db.myCollection.insert({someField: 'someValue'}) - to insert
    - db.myCollection.find() - to find all documents
    - db.myCollection.find({predicate}, {projection}) - to find all documents
      - projection: {field1: 1, field2: 0} ⇒ \_id will be included
    - db.myCollection.find({\_id: ObjectId("asdfasdfasdf")}) - to retrieve by ID
    - db.myCollection.find({\$and : [{someField: 'someValue'}, {predicate2}, {predicate3}]})
    - db.myCollection.find({someField: 'someValue'})
    - db.myCollection.update()
      - db.myCollection.update({predicate}, {the update})
      - db.myCollection.update({predicate}, {\$set: {field1: newValue1, field2: newValue2}})
      - db.myCollection.update({predicate}, {field1: newValue1, field2: newValue2}) – this replaces the document
    - db.myCollection.remove()
  - Mongoose
    - schemas
      - const mySchema = mongoose.Schema({
        - title: String,
        - author: {type: String, required: true},
        - semester: {type: String, enums: ["FALL", "SPRING", "SUMMER"]}
        - published: Date,
        - seats: {type: Number, default: 24}
      - }, {collection: "courses"})
    - models
      - const myModel = mongoose.model("MyModel", mySchema)
      - myModel.find()
      - myModel.findById()
      - myModel.findOne()
      - myModel.create()
      - myModel.updateOne({WHERE}, {UPDATE})
      - myModel.removeOne()
  - SQL Injection

■ ~~Types:~~

● ~~1 = 1~~

● ~~OR " = "~~

■ **PreparedStatement()**

- ps = connection.prepareStatement("SELECT \* FROM myTable WHERE someField=?")
- ps.setString(1, "someValue")

○ Transactions:

■ ACID

- Atomic
- Consistent
- Isolate
- Durability - permanent

■ Demarcate transactions

- turn off auto commit
- commit - make changes permanent
- rollback - undo any changes in the transaction

■ Potential problems

- Dirty Read
- Non-Repeatable reads
- Phantom reads

Isolation levels

READ-COMMITTED

REPEATABLE-READS

SERIALIZABLE

slowest

○ Indexes

■ Full table scan

■ Unique index on a primary key

■ Non-unique index on a non primary key

■ Optimize File IO

- buffers containing multiple records

■ Given a set of realistic records and File IO reads, compute whether its better to use index or full scan

■ Composite keys

○ ~~XML~~

■ ~~redundancy~~

■ ~~inconsistency~~